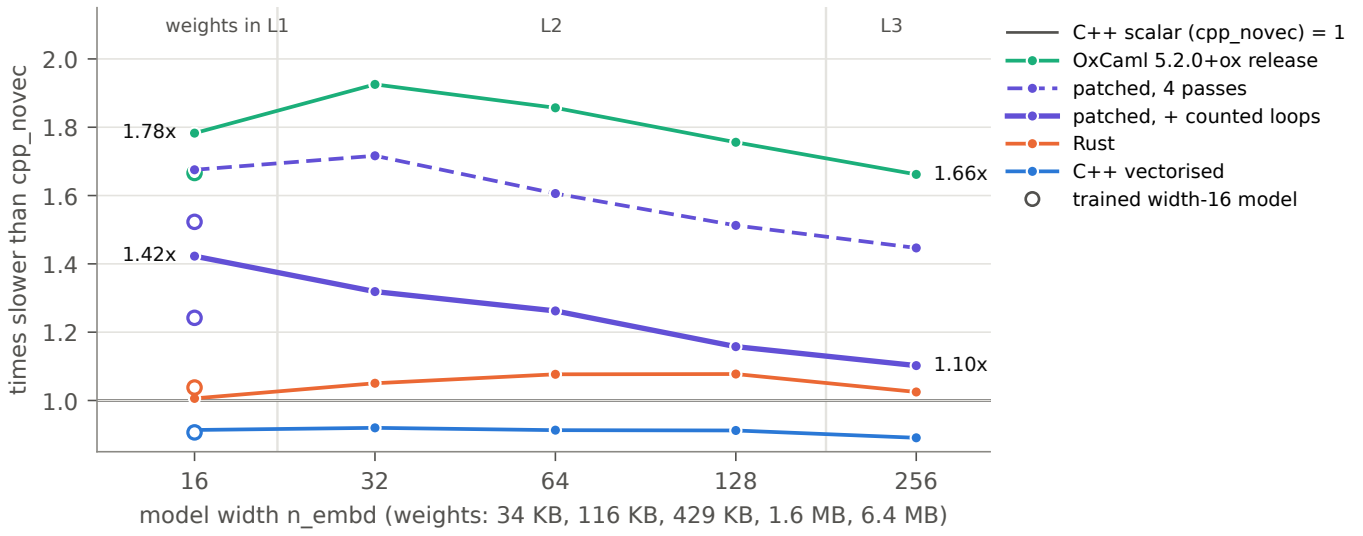
weights in L1
L2
L3
C++ scalar (cpp_novec) = 1
OxCaml 5.2.0+ox release
patched, 4 passes
patched, + counted loops
Rust
C++ vectorised
trained width-16 model
1.78x
1.66x
1.42x
1.10x
times slower than cpp_novec
2.0
1.8
1.6
1.4
1.2
1.0
16
32
64
128
256
model width n_embd (weights: 34 KB, 116 KB, 429 KB, 1.6 MB, 6.4 MB)